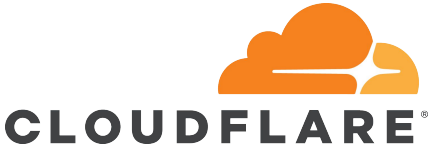

Flows and BGP, a multi-Terabyte story

Louis Poinsignon - Cloudflare

Why monitoring

- Billing
  - Reducing costs
- Traffic engineering
  - Where should we peer?
  - Where should we set-up a new PoP?
  - Optimizing our network
- Anomaly detection
  - Troubleshooting
  - Proactive monitoring and predictions

Ways to monitor

- SNMP
- Flows
- BGP

What this presentation will be about

CLOUDFLARE

Flows

The previous pipeline

- Single machine
  - High CPU usage
  - Not easily accessible
- Dropping packets (around 10%)
- Not monitored
- sFlow cannot be aggregated with NetFlow
- Data not accessible to other teams / hard to develop on

- nfacctd + opentsdb

Time for an upgrade

- Let’s use the tools provided by the platform team
  - Load-balancers
  - Mesos
  - Containers
  - Databases
  - Messages brokers
  - Big data processing clusters

CLOUDFLARE

What we built

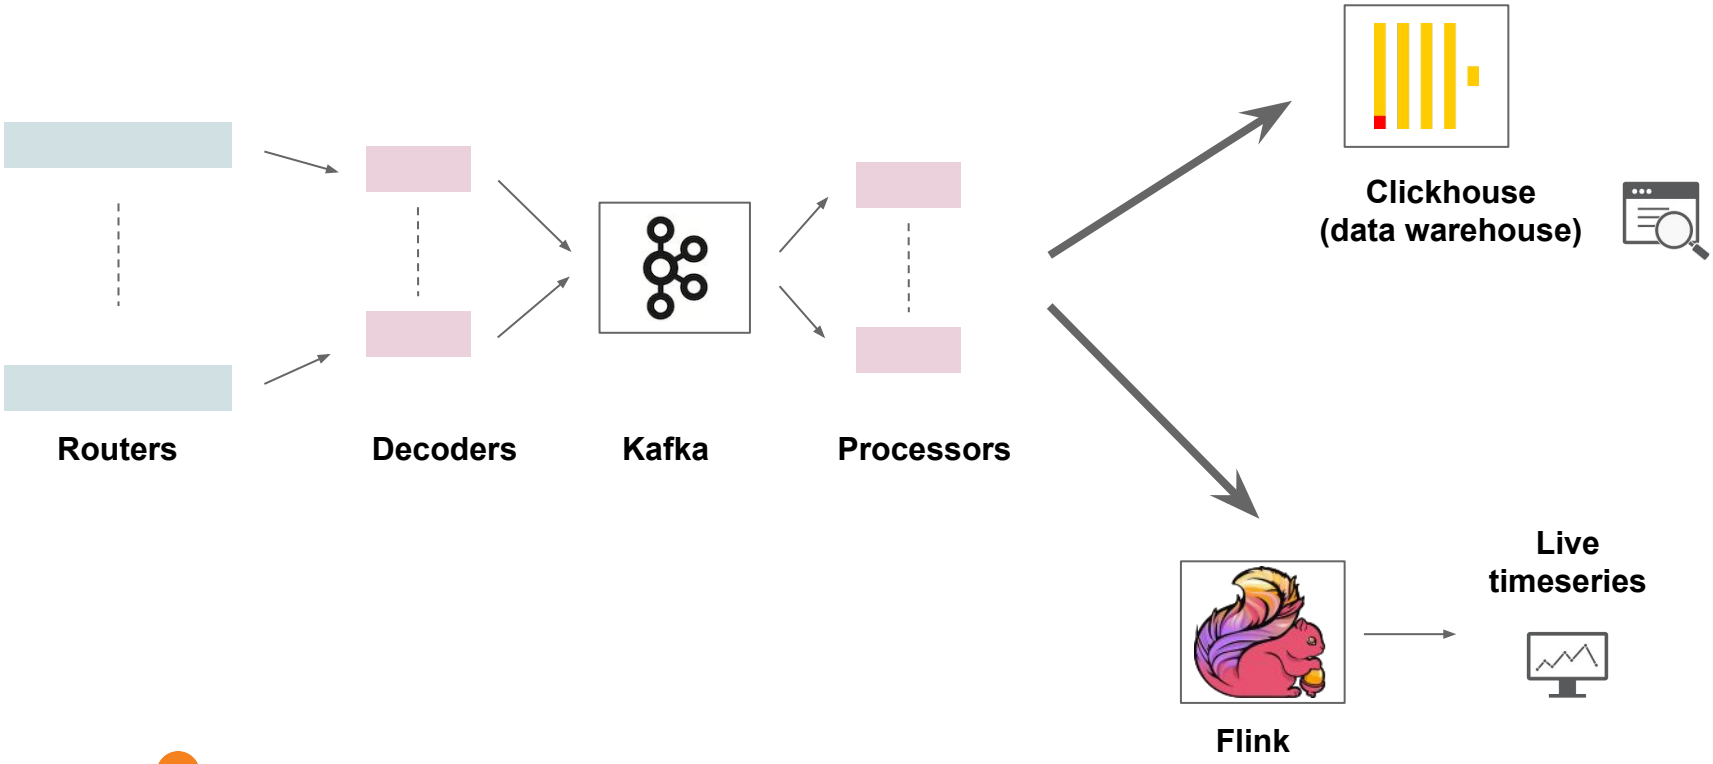

CLOUDFLARE

Flink - MapReduce

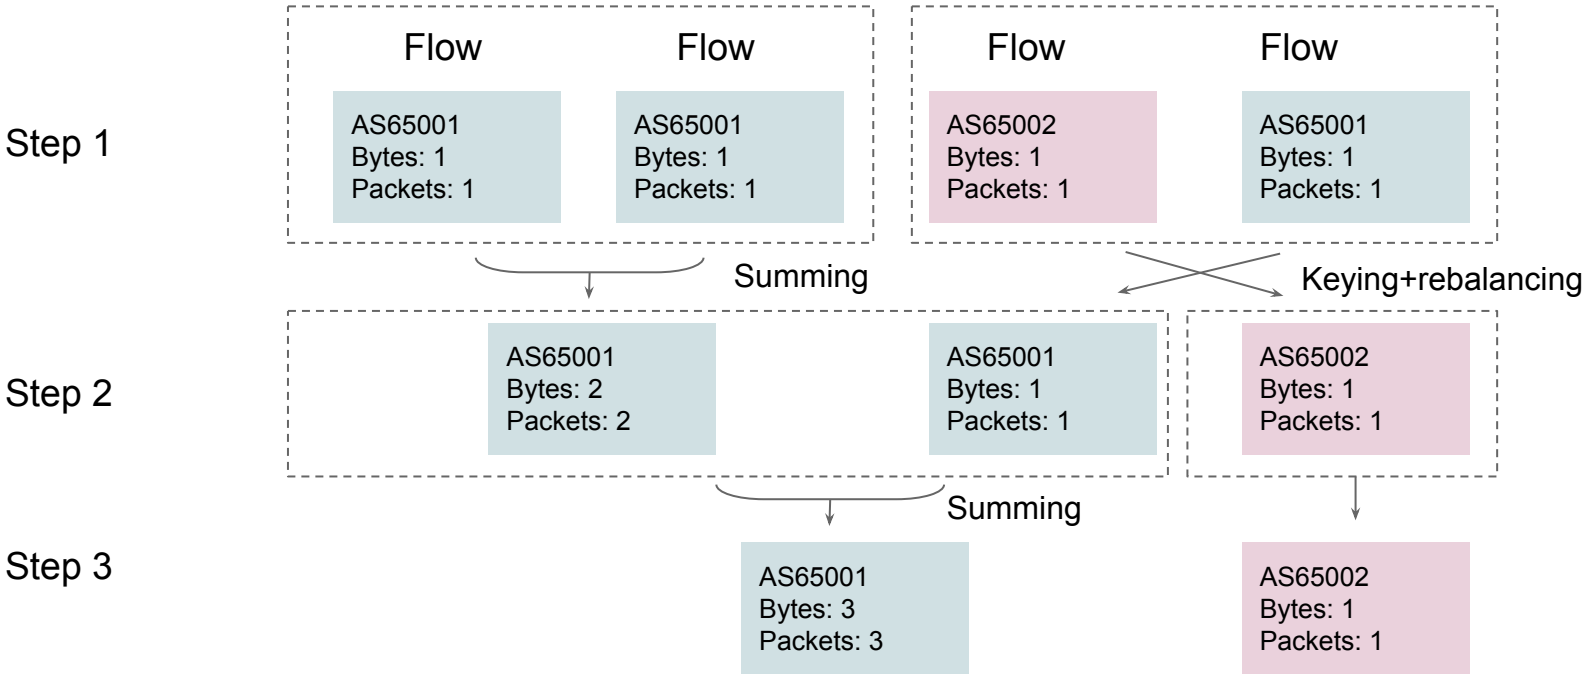

CLOUDFLARE

Flink - Sample program

```
DataStream<FlowMessage> inData =
new FlinkKafkaConsumer09<FlowMessage>(
                "netflows-processed",
                new FlowMessageDeserializer(),
                propertiesConsumer);

DataStream<FlowMessage> inDataEyeball =
                inData.filter(new FlowFilter.EyeballFilter()).
                setParallelism(1).broadcast();

DataStream<FlowAggMessage> inDataAgg =
                inDataEyeball.map(new FlowUtils.Mapper("DstAS,colo"));

inDataAgg.reduce(new FlowTransformations.FlowAggReduceKey());
```

Source (Kafka)

Filter

Mapping

Reduce

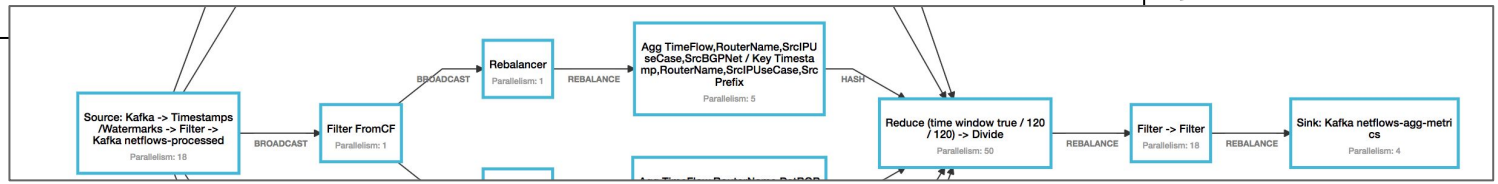

CLOUDFLARE

Flink - Windowing

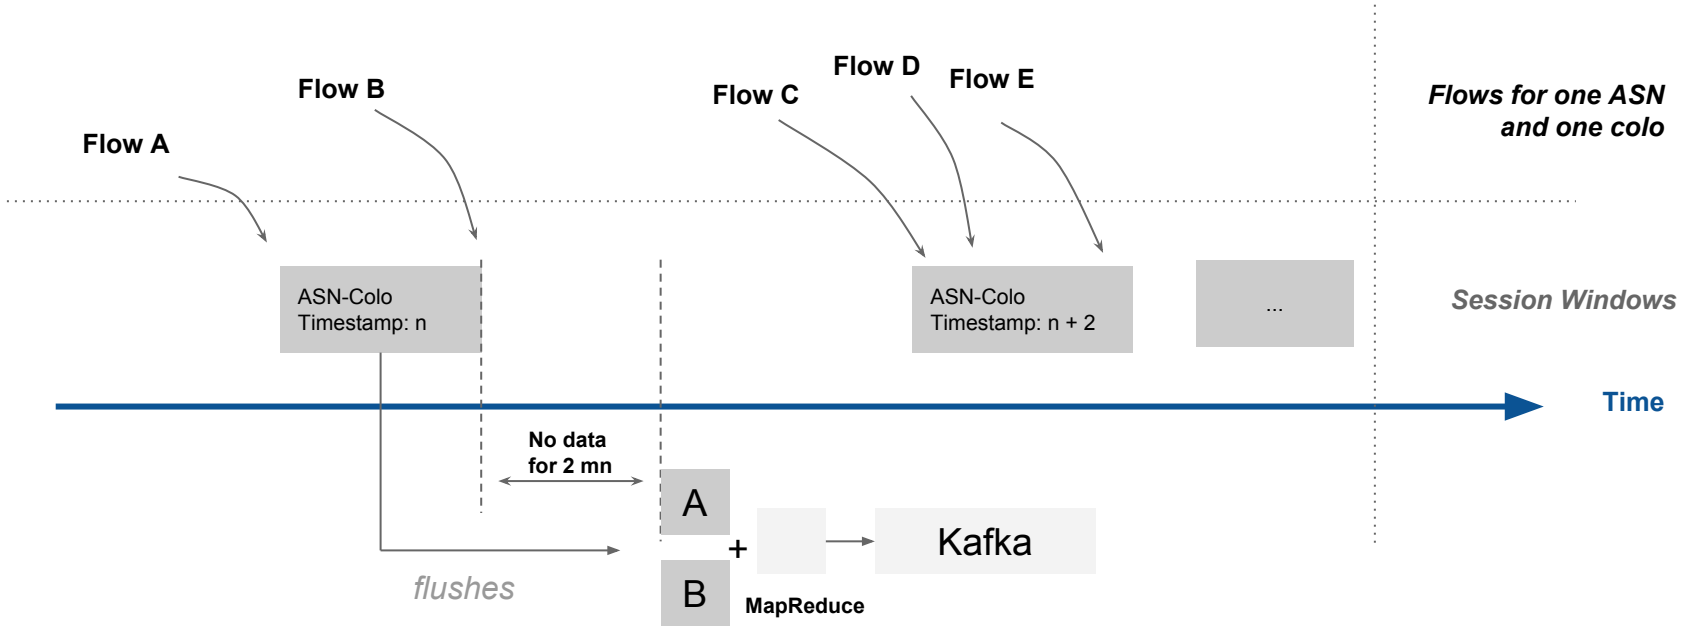

CLOUDFLARE

Results - Aggregations

Traffic generated by all the datacenters, by country, by network, etc.

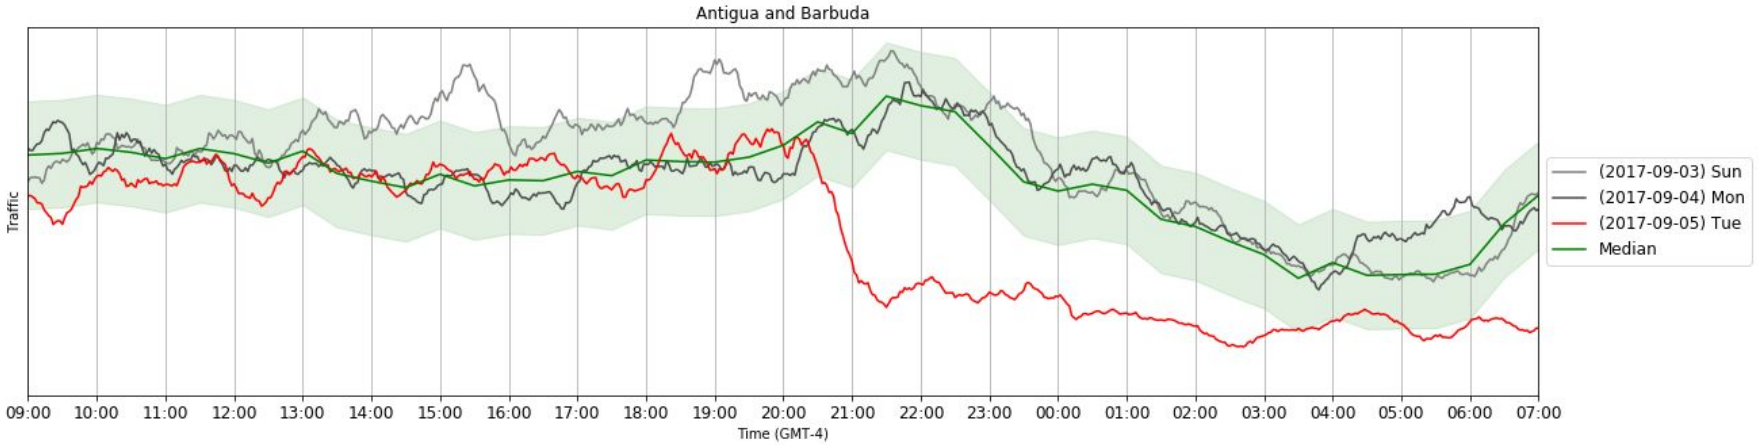

CLOUDFLARE

https://blog.cloudflare.com/the-story-of-two-outages/

Results - Statistics

Top providers per country, IPv6 penetration, etc.

	Network	AS	Ratio IPv6
1	Virgin Media	5089	0%
2	BTnet (BT's UK IP Network - AS2856)	2856	27.92%
3	Sky Broadband	5607	69.77%
4	EE	12576	5.19%

CLOUDFLARE

Results - Example: maintenance

Building a list with the best hours for maintenance

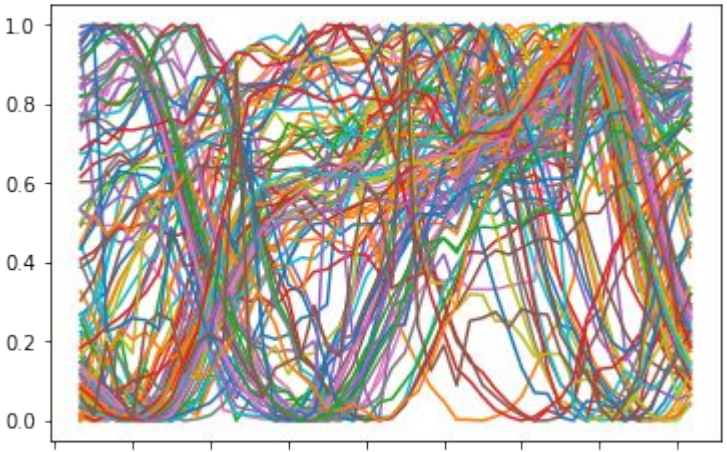

Normalized traffic of our PoP over a day

CLOUDFLARE

Algorithms

Derivation

Correlation

Pearson coefficient

Quantify the difference

Median

Remove small variations

Variance

Variation intensity

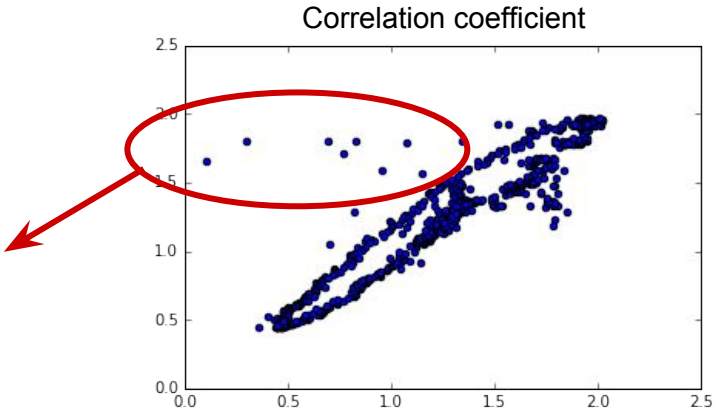

Outliers

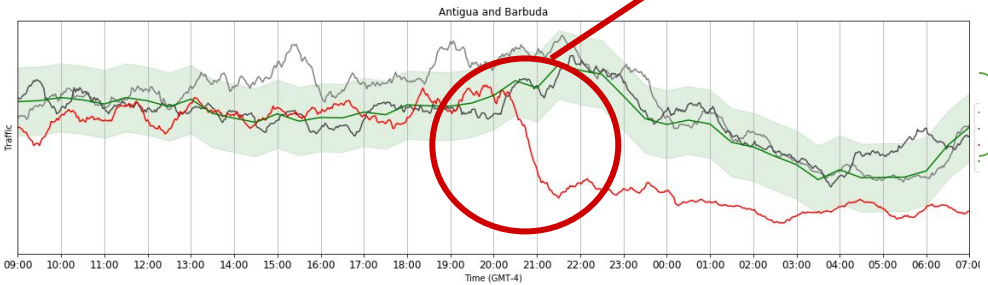

Local variance following median

CLOUDFLARE

BGP

BGP flaps

Hi,

Your PeeringDB entry suggests 15000, but we have the maximum set to 20,000, hoe many prefixes are you sending?

Maybe we announce more than 20,000. At the moment I can not say for sure.
Please, increase more.

I can not say for sure.

Please, increase more.

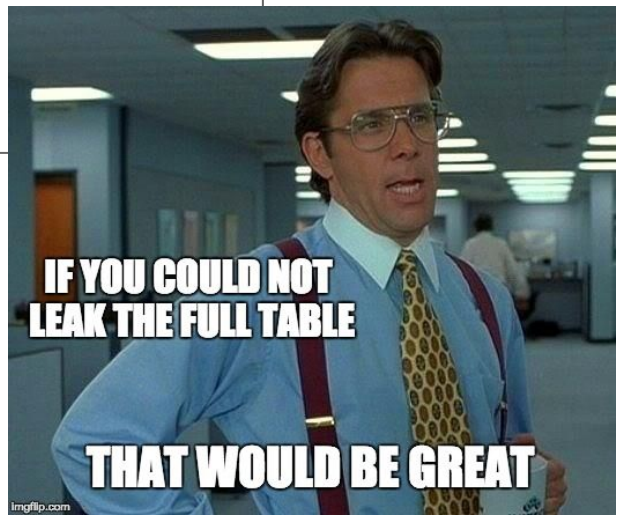

CLOUDFLARE

BGP

100+ routers with full tables

Want to be aware of route-leaks or other anomalies

Need to do data analysis and periodic reports

Provide improved data for the flow pipeline

- Mappings prefixes to ASN
- Next-hop to ASN
- Peering/transit information

BGP Pipeline

Completely custom BGP implementation to live on a cloud with elastic IPs.

- More flexible than GoBGP

Horizontal scaling and hashing

- A single machine contains a subset of the data.
- Messages from x will always reach y

Develop custom trie implementation

- Stores IP prefixes in an optimized way

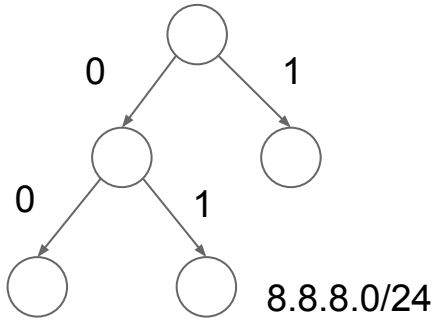

CLOUDFLARE

BGP Pipeline

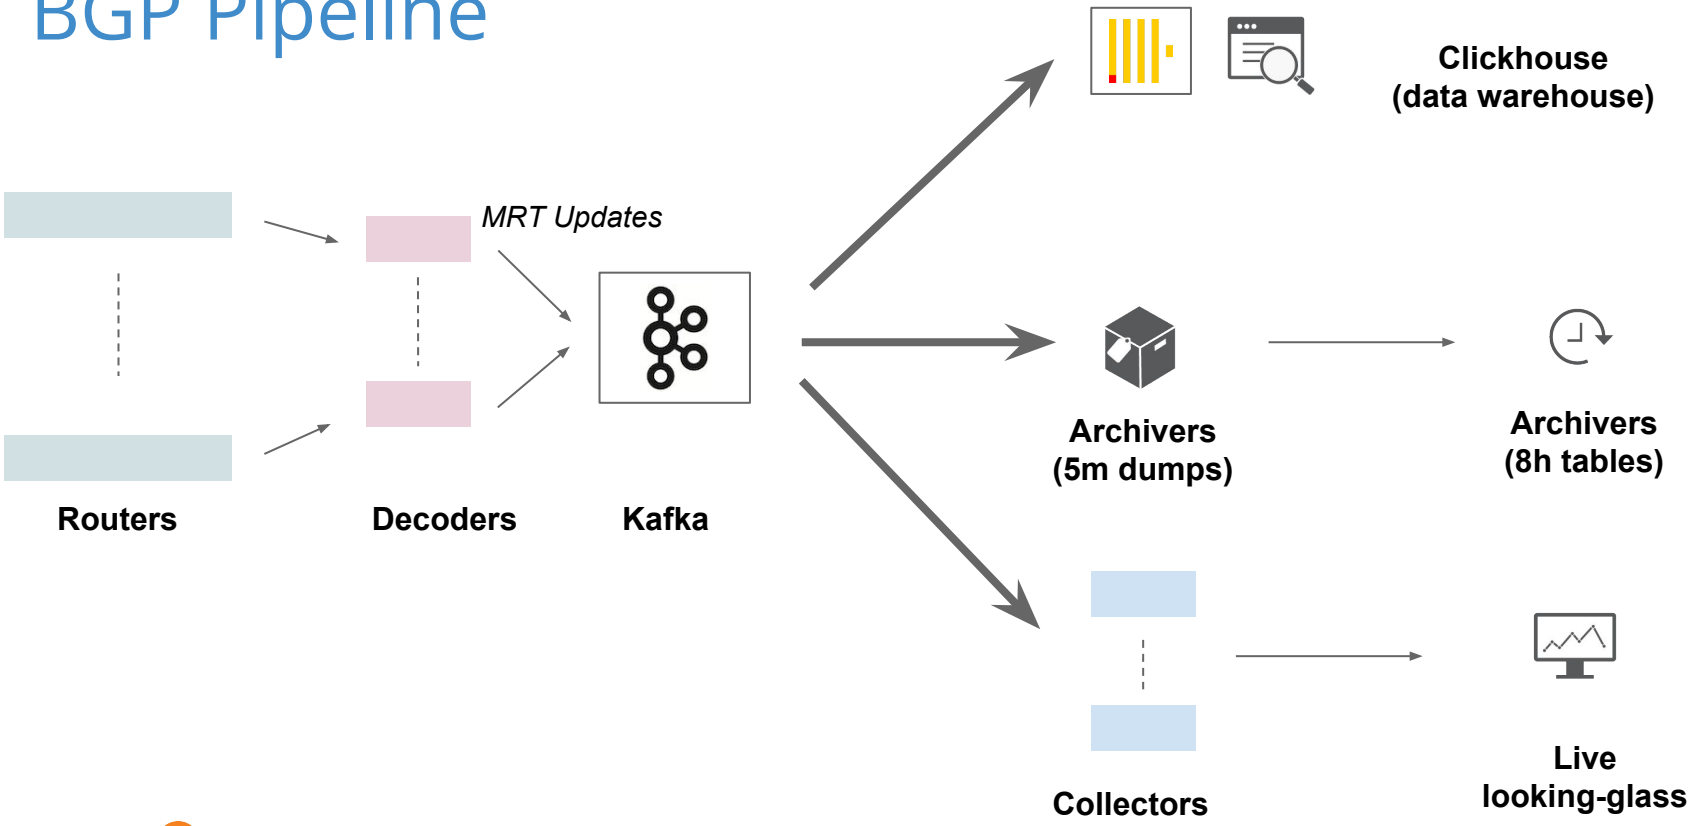

CLOUDFLARE

BGP API

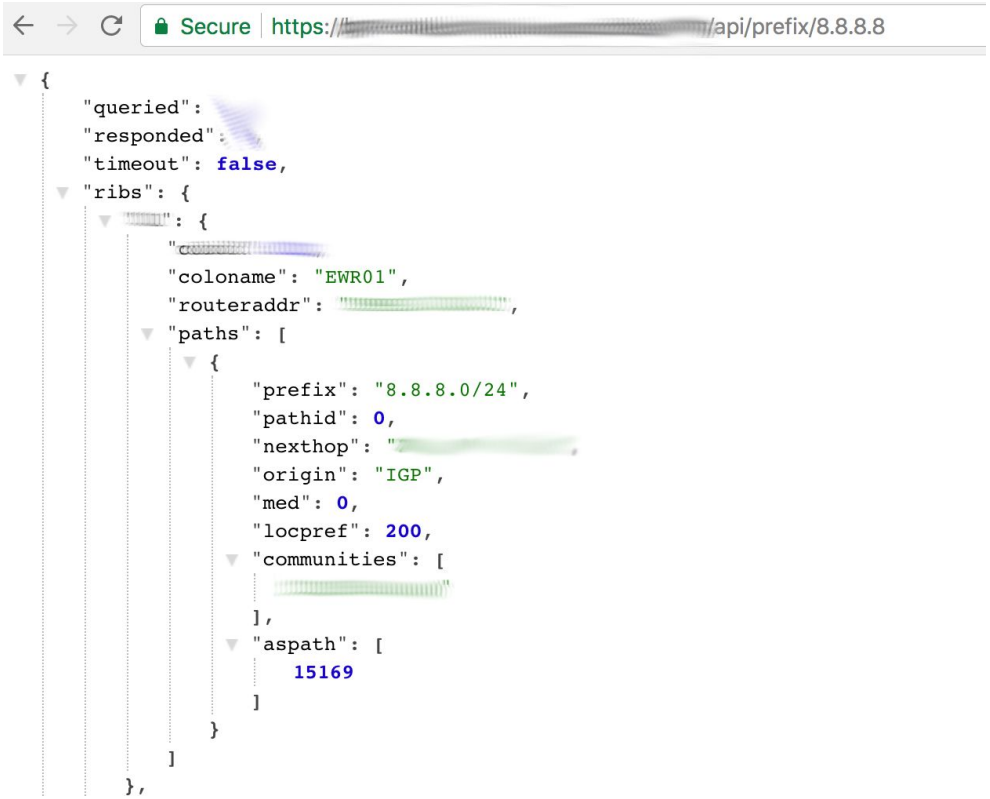

CLOUDFLARE

Discoveries

- People sending us IX LAN prefixes
- Receiving smaller than /48 IPv6 and smaller than /24 IPv4
- Longest AS-Path
  - 2402:8100:3980::/42 → 37 ASNs

CLOUDFLARE

A train

"We have approximately 170 000 routes"

CLOUDFLARE

TATA’s Peeringdb

IPv4 Prefixes	220000
IPv6 Prefixes	11000

CLOUDFLARE

Transit

3 types of clients. No BMP available.

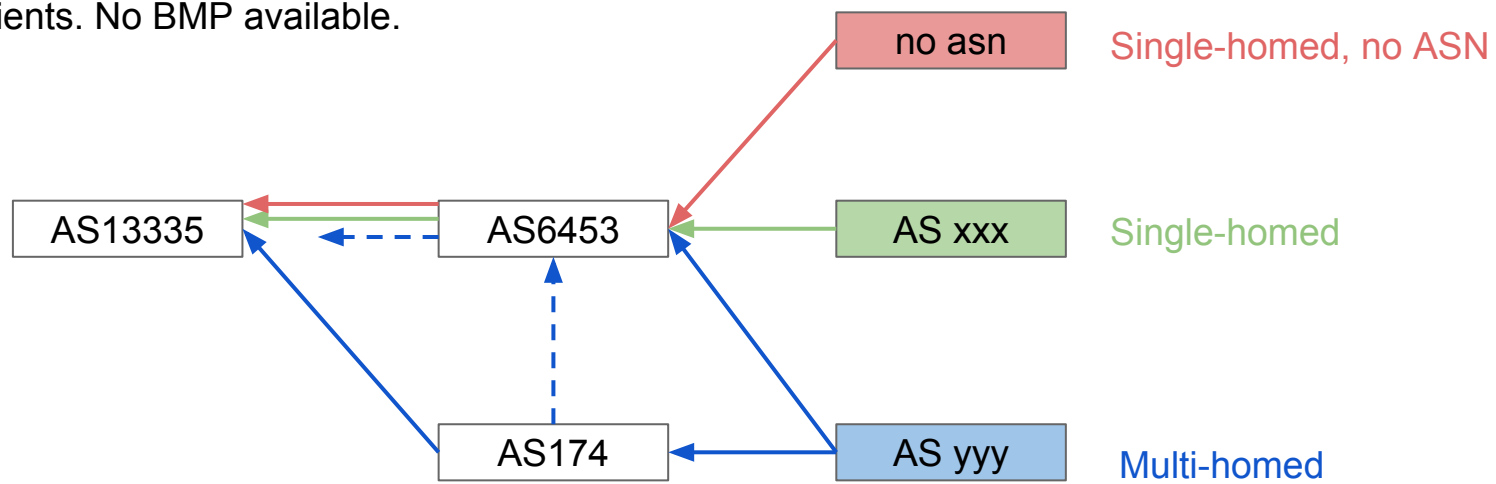

CLOUDFLARE

Let’s process

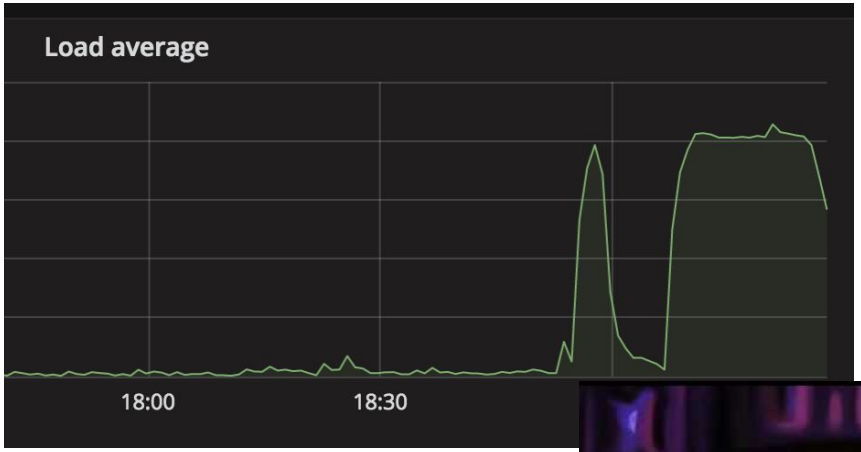

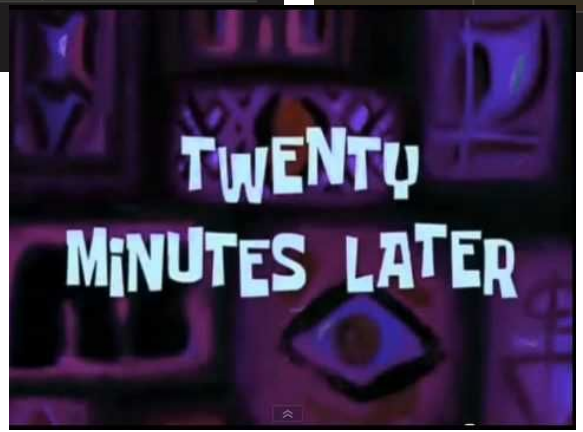

CLOUDFLARE

Results

315997 routes (approx 42%)

- with 6453 as the last ASN with no transit listed behind

205 routes

- With 6453 as the last ASN

Around 100k routes per PoP

- With 6453 as ASN

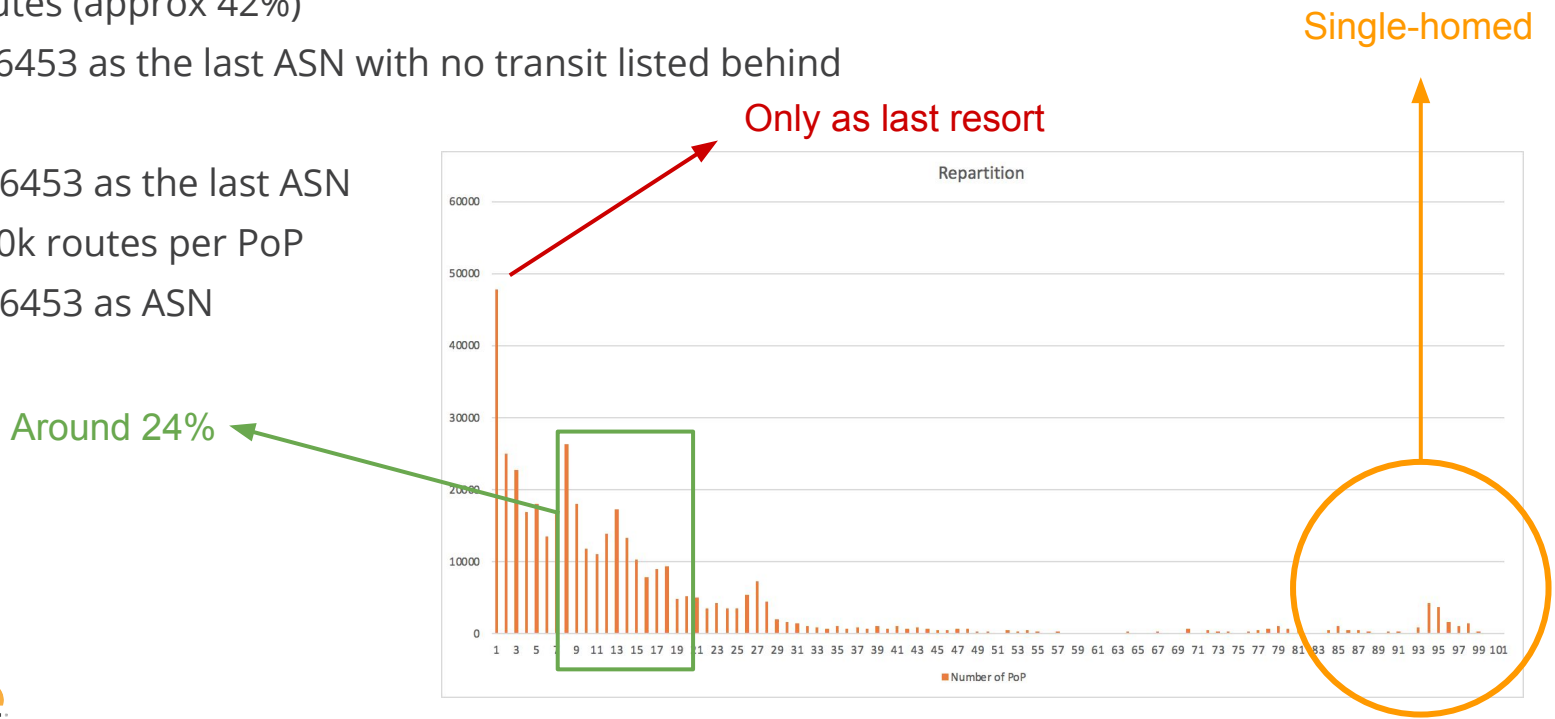

CLOUDFLARE

IP prefix sizes

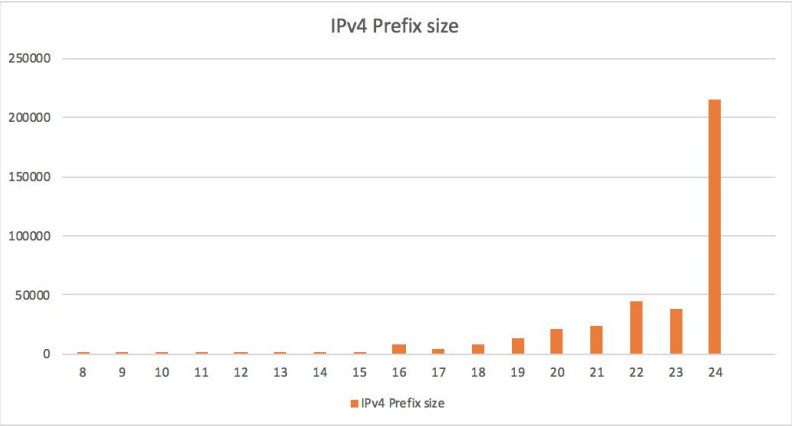

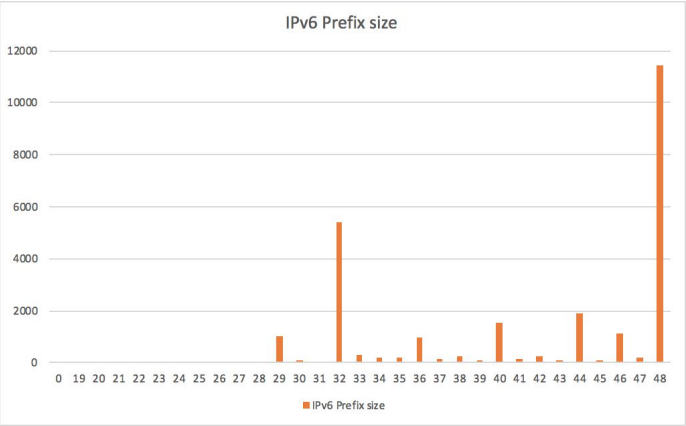

CLOUDFLARE

Size, costs

Size

Flows

- Raw feed: 10-50 mbps
- Unique storage: 1-5 Terabytes
- Aggregated storage: 10-100 GigaBytes

BGP

- Raw feed: few kbps (100s updates per second)
- Around 35 GB of RAM used for storing the tables
- Around 100 MB per full table stored in MRT format

Costs

The price of running it on a cloud (Amazon/Google/Azure) will depend on your size.

Rule of thumb:

traffic	10Gb/s	100Gb/s	1Tb/s
routers	10	50	100
Flows	$200/mo	$800/mo	$1,500/mo
BGP	$200/mo	$600/mo	$800/mo

CLOUDFLARE

Open-source

Open-source (soon)

Flows:

- Decoders
- ~~Processors~~ → too specific to Cloudflare
- Flink example

BGP:

- Library to manage a session
- ~~Server, APIs~~ → too specific to Cloudflare

Tools:

- Maxmind DB encoder

CLOUDFLARE

Thank you!
Questions?